

Q1. Product Reviews 1a) Target (continuous) Future sales volume (units) over a fixed horizon (e.g., next 30 days). 1b) Linear model with rating + judgment words Let Score be the numeric rating. Let freq_good, freq_bad, freq_doesnt_work be normalized word frequencies (e.g., per 100 tokens). Model: $\text{SalesNext30d} = b_0 + b_1 \cdot \text{Score} + b_2 \cdot \text{freq_good} + b_3 \cdot \text{freq_bad} + b_4 \cdot \text{freq_doesnt_work} + \text{error}$ 1c) Reducing features - LASSO (L1) regularization to shrink unhelpful coefficients to zero. - Dimensionality reduction (e.g., PCA on text features). - Aggregate to a single sentiment score instead of multiple word counts. 1d) Mixed scales (1-5 vs 1-10) Normalize ratings to a common scale before modeling (e.g., convert 1-10 to 1-5 by dividing by 2, or standardize to z-scores). Optionally add an indicator feature for original scale (e.g., is_scale_10) to allow different intercepts/slopes if needed. Without normalization, the Score coefficient is not comparable across scales. Q2. Fruit dataset 2a) Categorical variables type, color, size 2b) Total variables after one-hot coding - If you one-hot encode all levels (no drop): 7 dummy vars + price = 8 total. - If you use drop-first (to avoid multicollinearity with an intercept): 4 dummy vars + price = 5 total. (Counting only the encoded features: 7 vs 4.) Q3. Small linear regression Given: House 1: Size=1400, Beds=3, Price=245 House 2: Size=1600, Beds=3, Price=312 House 3: Size=1700, Beds=4, Price=279 (Price in 1000s) Model: $\text{Price} = \beta_0 + \beta_1 \cdot \text{Size} + \beta_2 \cdot \text{Bedrooms}$ (We'll scale Size by 100 for easier arithmetic: $\text{Size}_{100} = \text{Size}/100$) 3a) Feature matrix A
$$= \begin{bmatrix} 14 & 3 \\ 16 & 3 \\ 17 & 4 \end{bmatrix}$$
 3b) Target vector $y = [245, 312, 279]^T$ 3c) Normal equation $\beta = (A^T A)^{-1} A^T y$ gives $\beta = [\beta_0, \beta_1, \beta_2]$
$$= [-24.5, 33.5, -66.5]$$
 (If you do NOT scale Size: $\beta = [-24.5, 0.335, -66.5]$) 3d) Predict for Size = 1500 sqft ($\text{Size}_{100} = 15$), Bedrooms = 3: $\text{Price} = -24.5 + 33 \cdot 15 + 278.5 = 278.5$ (in 1000s) \Rightarrow Predicted price \approx \$278,500 Q4. When is $(A^T A)$ invertible? Claim. For $A \in \mathbb{R}^{n \times p}$, $(A^T A)$ is invertible \Leftrightarrow the columns of A are linearly independent. Proof. (\Rightarrow) Suppose $A^T A$ is invertible. If A 's columns were dependent, there would exist $x \neq 0$ with $Ax = 0$ (i.e., a nontrivial vector in the null space of A). Then $A^T A x = A^T (Ax) = A^T 0 = 0$, so x is in the null space of $A^T A$. But an invertible matrix has only the zero vector in its null space. Contradiction. Hence A 's columns are linearly independent. (\Leftarrow) Suppose A 's columns are linearly independent. Then $\text{rank}(A) = p$ (full column rank), so for any $x \neq 0$ we have $Ax \neq 0$ and $x^T A^T A x = \|Ax\|^2 > 0$. Thus $A^T A$ is symmetric positive definite and therefore invertible. (Equivalently, by the rank identity $\text{rank}(A^T A) = \text{rank}(A) = p$; since $A^T A$ is $p \times p$ with rank p , it is invertible.)

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_california_housing
from sklearn.linear_model import LinearRegression

california = fetch_california_housing(as_frame=True)
X_df = california.data
y = california.target

print("5a) Feature names:", list(X_df.columns))

print("\n5b) Rows 10-15 for HouseAge, AveRooms, Population:")
print(X_df.loc[10:15, ['HouseAge', 'AveRooms', 'Population']])

# 5c) Simple linear regression: MEDV ~ AveOccup
X = X_df[['AveOccup']].values
n = len(y)
split = (2*n)//3
X_train, X_test = X[:split], X[split:]
y_train, y_test = y.values[:split], y.values[split:]

model = LinearRegression()
model.fit(X_train, y_train)

print("\n5d) R^2 (train):", model.score(X_train, y_train))
print("5d) R^2 (test):", model.score(X_test, y_test))
print("Intercept:", model.intercept_)
print("Slope for AveOccup:", model.coef_[0])

plt.figure()
sample = np.arange(0, n, 50)
plt.scatter(X_df[['AveOccup']].iloc[sample], y.iloc[sample], s=8)
xs = np.linspace(X.min(), X.max(), 200).reshape(-1,1)
ys = model.predict(xs)
plt.plot(xs, ys, linewidth=2)
plt.xlabel("AveOccup")
plt.ylabel("MedHouseValue ($100k)")
plt.title("Simple Linear Regression: MedHouseValue ~ AveOccup")
plt.show()
```

5a) Feature names: ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']

5b) Rows 10-15 for HouseAge, AveRooms, Population:

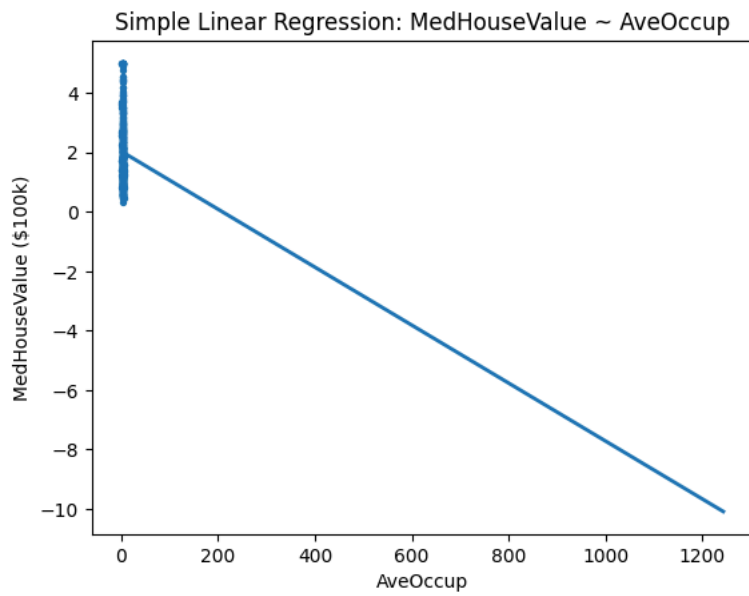
| | HouseAge | AveRooms | Population |
|----|----------|----------|------------|
| 10 | 52.0 | 5.477612 | 910.0 |
| 11 | 52.0 | 4.772480 | 1504.0 |
| 12 | 52.0 | 5.322650 | 1098.0 |
| 13 | 52.0 | 4.000000 | 345.0 |
| 14 | 52.0 | 4.262903 | 1212.0 |
| 15 | 50.0 | 4.242424 | 697.0 |

5d) R^2 (train): 0.0023374510886858824

5d) R^2 (test): -0.03700899339889241

Intercept: 2.038064223107151

Slope for AveOccup: -0.009751795964041849



5e) R^2 : fraction of target variance explained; expect low with one predictor. Big train-test gap \Rightarrow overfitting; both low \Rightarrow weak model. Q6: 6a) $b \notin \text{col}(A)$ 6b) $b \in \text{col}(A)$ and $\text{rank}(A)=n$ 6c) $b \in \text{col}(A)$ and $\text{rank}(A)<n$